Parallel Image Processing with OpenMP: Greyscale, Normalisation and Resize

MTCS-103P — High Performance Computing Laboratory · Mini Project P1: Image Processing (2), Part 2

Submitted by: M. Sai Praneeth — Registration no. 26010303004
Implementation: C99 + OpenMP
Repository: https://github.com/buvann-11/Parallel-Image-Processing-OpenMP

1. Objective

Before an image reaches a neural network it is usually converted to greyscale, **normalised** to a standard range and **resized** to a fixed input size. Each step touches every pixel, so on large images they are worth parallelising.

This work implements the three operations in C, parallelises each with OpenMP, verifies the results exactly, measures how well each one scales and explains why some scale better than others. It also provides a GUI for running the operations and the test suite.

2. Results at a glance

Kernel	Serial (ms)	2 threads (ms)	Speedup
Greyscale, gamma-correct luminance	36.12	18.12	1.99×
Greyscale, BT.601	0.60	0.32	1.87×
Normalise, min–max	5.95	2.77	2.14×
Normalise, histogram equalisation	6.57	4.29	1.53×
Resize, bicubic	16.91	10.80	1.57×



3. Algorithms

3.1 Greyscale

A weighted sum of red, green and blue, because the eye is most sensitive to green: BT.601 (0.299, 0.587, 0.114), BT.709 (0.2126, 0.7152, 0.0722) and the plain mean. A fourth method computes true **linear luminance**: sRGB values are gamma-encoded (128 is about 21 % of the light of 255, not 50 %), so pixels are linearised, weighted and re-encoded — two `powf` calls per pixel.

3.2 Normalisation

Every method needs a **global** statistic before any output pixel can be written, so each runs as a parallel reduction followed by a parallel map:

Method	Phase 1 — reduce	Phase 2 — map
min–max	global min and max	rescale to 0–255
z-score	Σx and Σx^2 in one pass	rescale to mean 128, $\pm 2\sigma$
equalisation	256-bin histogram	lookup table from the cumulative histogram

3.3 Resize

Each output pixel is a weighted average of nearby input pixels (nearest, bilinear, or bicubic with Keys $a = -0.5$), sampled at pixel centres as in Pillow and OpenCV. Two design points:

- **Separable, two passes.** Horizontal then vertical, so a bicubic resize does 4 + 4 taps per pixel instead of 4 × 4 — 4× less arithmetic before any parallelism.
- **Anti-aliasing.** When shrinking, the kernel is widened by the scale factor so every input pixel contributes; otherwise fine detail turns into moiré.

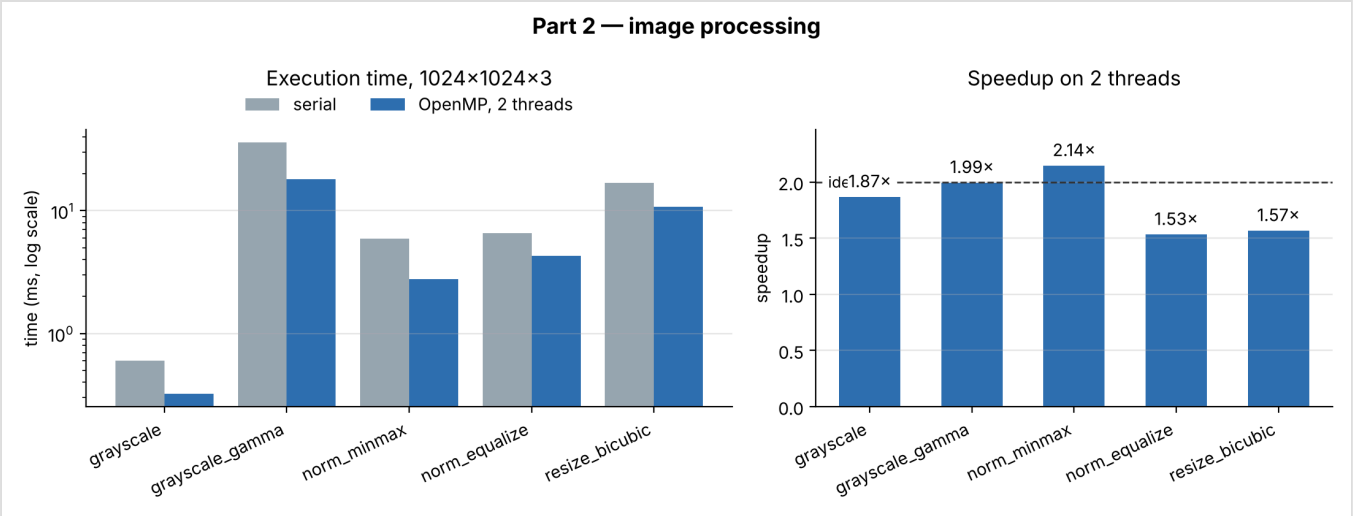
4. Parallelisation

Every kernel has a `_serial()` version with no OpenMP and an `_omp()` version; speedup is measured against the directive-free one.

- **Greyscale** — one `parallel for` over pixels; each thread writes only its own pixels.
- **Min–max** — `reduction(min:lo) reduction(max:hi)`.
- **Z-score** — two `reduction(+:)` sums in a single pass over memory.
- **Equalisation** — each thread fills a **private histogram** padded to whole 64-byte cache lines; the copies are merged afterwards. A shared histogram with `#pragma omp atomic` makes every thread fight over the same cache lines and runs slower than serial.
- **Resize** — the horizontal and vertical passes are separate parallel loops; the implicit barrier between them is the only synchronisation needed. The small weight tables are deliberately built serially.

5. Performance

Strong scaling on a 1024×1024×3 input, best of repeated runs after a warm-up, 2-core x86-64 Linux machine, GCC `-O3`.



Arithmetic intensity predicts scalability. Gamma-correct greyscale does expensive arithmetic per pixel and reaches the ideal 1.99×. Plain greyscale does one multiply-add per byte loaded; it is memory-bound and reaches 1.87×.

Serial phases cap the rest. Equalisation (1.53×) has a serial histogram merge and table build; resize (1.57×) has serial weight tables and a barrier. The Karp–Flatt metric, `e = (1/S − 1/p)/(1 − 1/p)`, measures the serial fraction directly: $e \approx 0.31$ for equalisation and 0.28 for resize, matching those serial sections in the code.

Chunk size is not a detail. With `schedule(dynamic)` OpenMP's default chunk is 1, so every pixel passes through the runtime's dispatcher:

Kernel	**`static`**	**`dynamic`, chunk 1**	**`dynamic`, chunk 8192**
Greyscale	1.32×	**0.01×**	1.37×
Normalise, min–max	2.18×	**0.02×**	2.02×
Normalise, equalisation	1.50×	**0.02×**	0.97×
Resize, bicubic	1.68×	1.63×	**0.90×**

Per-pixel loops become 100× *slower* at chunk 1; resize, whose iterations are whole rows, instead starves at chunk 8192. The chunk must match the loop granularity, and for uniform work `static` is the right choice.

Baselines matter. An early version reported 15× on 2 cores. The cause: the "serial" baseline was `#pragma omp parallel for if(0)`, and GCC still runs the OpenMP loop machinery when the condition is false. Directive-free serial code fixed the measurement.

6. Correctness

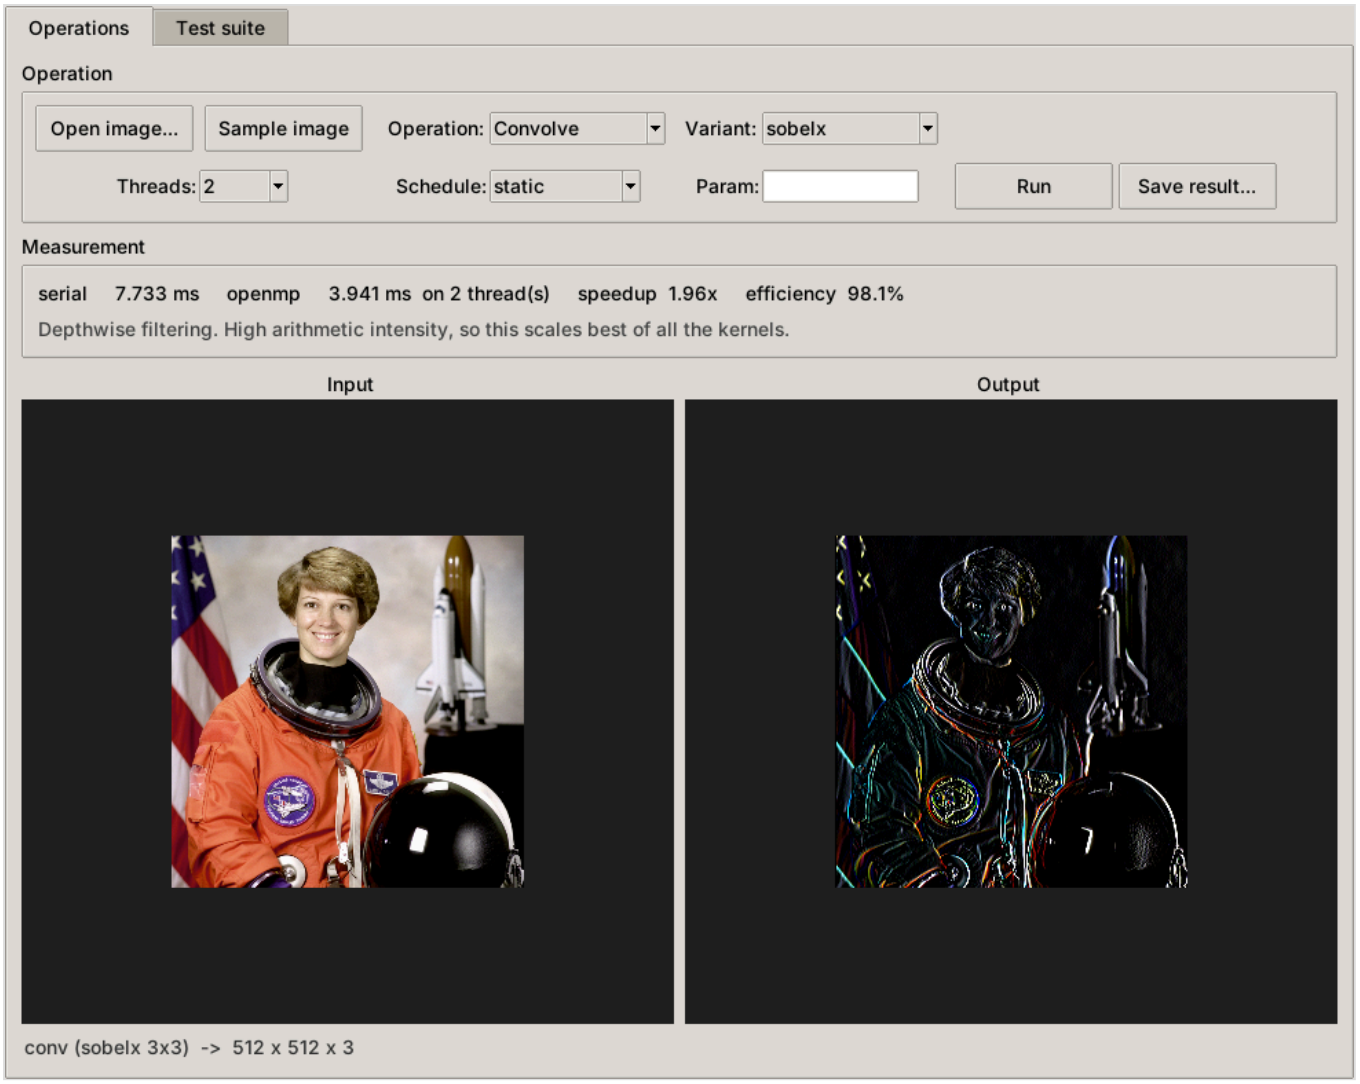
Check	Result
OpenMP output = serial output	exact, every method
Output at 1–4 threads; static, dynamic and guided	bit-identical
Independent NumPy implementation	**0** grey-level difference for greyscale, min–max, equalisation and nearest resize; ±1 for bilinear and bicubic
ThreadSanitizer with Archer; AddressSanitizer	no races; clean
Linux, macOS and Windows on every push	pass

Two real bugs were found and fixed:

- **Resize.** Nearest-neighbour downscaling disagreed with NumPy on 0.6 % of pixels: in `float32`, `floor((p + 0.5) · scale)` landed on the wrong side of an integer for ratios like 256/97. Sample positions are now computed in `double`.
- **Compiler bug in equalisation.** On macOS, Apple clang 21 at `-O2`/`-O3` dropped counts in the per-thread histogram loop under `schedule(runtime)` — 27,648 pixels gave only 19,633 counts — so equalisation differed from the serial version. GCC and `-O0`/`-O1` were correct. Writing each increment through a `volatile` pointer avoids the faulty optimisation at no measurable cost.

7. Graphical interface

`gui/app.py` uses only the Python standard library. The **Operations** tab loads any image (PPM, PNG, JPEG, BMP, TGA), runs an operation with a chosen thread count and schedule, and shows input, output, serial time, parallel time and speedup side by side.



The **Test suite** tab runs each test stage with one click — C test suite, image-format tests, NumPy cross-check, a thread-count sweep that hashes every output, and a quick benchmark — with a live PASS/FAIL badge for each.

8. Conclusions

1. **Arithmetic intensity decides scaling** — expensive per-pixel arithmetic scales ideally, memory-bound maps do not.
2. **Serial phases set the ceiling** (Amdahl), and Karp–Flatt measures them.
3. **Chunk size must match loop granularity** — a wrong choice cost 100×.
4. **Algorithms come first** — the separable resize saves 4× arithmetic, more than a second core provides.
5. **Exact serial = parallel testing finds real bugs**, including one in the compiler.